

## Importing Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [2]: df = pd.read_csv('/content/Customer_Data.csv')
```

## Data Cleaning and Exploration

```
In [3]: df.head()
```

```
Out[3]:
```

	CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLME
0	C10001	40.900749	0.818182	95.40	0.00	
1	C10002	3202.467416	0.909091	0.00	0.00	
2	C10003	2495.148862	1.000000	773.17	773.17	
3	C10004	1666.670542	0.636364	1499.00	1499.00	
4	C10005	817.714335	1.000000	16.00	16.00	

```
In [4]: # Removing English Letter from CUST_ID using regular expression

df['CUST_ID'] = df['CUST_ID'].str.replace(r'\D', '', regex = True)

# Change Data Type from object to int
df['CUST_ID'] = df['CUST_ID'].astype('int64')
```

```
In [5]: df.head()
```

```
Out[5]:
```

	CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	INSTALLME
0	10001	40.900749	0.818182	95.40	0.00	
1	10002	3202.467416	0.909091	0.00	0.00	
2	10003	2495.148862	1.000000	773.17	773.17	
3	10004	1666.670542	0.636364	1499.00	1499.00	
4	10005	817.714335	1.000000	16.00	16.00	

In [6]: *# Checking the Data Types of All Columns*

```
df.dtypes
```

Out[6]:

CUST_ID	int64
BALANCE	float64
BALANCE_FREQUENCY	float64
PURCHASES	float64
ONEOFF_PURCHASES	float64
INSTALLMENTS_PURCHASES	float64
CASH_ADVANCE	float64
PURCHASES_FREQUENCY	float64
ONEOFF_PURCHASES_FREQUENCY	float64
PURCHASES_INSTALLMENTS_FREQUENCY	float64
CASH_ADVANCE_FREQUENCY	float64
CASH_ADVANCE_TRX	int64
PURCHASES_TRX	int64
CREDIT_LIMIT	float64
PAYMENTS	float64
MINIMUM_PAYMENTS	float64
PRC_FULL_PAYMENT	float64
TENURE	int64
dtype:	object

In [7]: *# Change all the Data Types to float*

```
for col in df.columns:  
    if df[col].dtypes == 'int64':  
        df[col] = df[col].astype('float64')
```

In [8]: *# Information of the whole dataset*

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 8950 entries, 0 to 8949
```

```
Data columns (total 18 columns):
```

#	Column	Non-Null Count	Dtype
0	CUST_ID	8950 non-null	float64
1	BALANCE	8950 non-null	float64
2	BALANCE_FREQUENCY	8950 non-null	float64
3	PURCHASES	8950 non-null	float64
4	ONEOFF_PURCHASES	8950 non-null	float64
5	INSTALLMENTS_PURCHASES	8950 non-null	float64
6	CASH_ADVANCE	8950 non-null	float64
7	PURCHASES_FREQUENCY	8950 non-null	float64
8	ONEOFF_PURCHASES_FREQUENCY	8950 non-null	float64
9	PURCHASES_INSTALLMENTS_FREQUENCY	8950 non-null	float64
10	CASH_ADVANCE_FREQUENCY	8950 non-null	float64
11	CASH_ADVANCE_TRX	8950 non-null	float64
12	PURCHASES_TRX	8950 non-null	float64
13	CREDIT_LIMIT	8949 non-null	float64
14	PAYMENTS	8950 non-null	float64
15	MINIMUM_PAYMENTS	8637 non-null	float64
16	PRC_FULL_PAYMENT	8950 non-null	float64
17	TENURE	8950 non-null	float64

```
dtypes: float64(18)
```

```
memory usage: 1.2 MB
```

Checking the total missing values of each columns

```
In [9]: missing_values = []
        for col in df.columns:
            missing_values.append(df[col].isna().sum())
```

```
In [10]: Col = df.columns
```

```
In [11]: Col = pd.DataFrame(Col)
          missing_values= pd.DataFrame(missing_values)
```

```
In [12]: result_missing = pd.concat([Col, missing_values], axis = 1)
          result_missing.columns = ['Columns', 'Missing_values']
```

```
In [13]: result_missing
```

```
Out[13]:
```

	Columns	Missing_values
0	CUST_ID	0
1	BALANCE	0
2	BALANCE_FREQUENCY	0
3	PURCHASES	0
4	ONEOFF_PURCHASES	0
5	INSTALLMENTS_PURCHASES	0
6	CASH_ADVANCE	0
7	PURCHASES_FREQUENCY	0
8	ONEOFF_PURCHASES_FREQUENCY	0
9	PURCHASES_INSTALLMENTS_FREQUENCY	0
10	CASH_ADVANCE_FREQUENCY	0
11	CASH_ADVANCE_TRX	0
12	PURCHASES_TRX	0
13	CREDIT_LIMIT	1
14	PAYMENTS	0
15	MINIMUM_PAYMENTS	313
16	PRC_FULL_PAYMENT	0
17	TENURE	0

### Filling Missing Values


```
In [14]: df['CREDIT_LIMIT'].fillna(df['CREDIT_LIMIT'].mean(), inplace=True)
```

```
In [15]: from sklearn.impute import SimpleImputer
imputer = SimpleImputer(missing_values = np.nan, strategy = 'mean')
df[['MINIMUM_PAYMENTS']] = imputer.fit_transform(df[['MINIMUM_PAYMENTS']])
```

## Outlier Detection using Boxplot

```
In [16]: plt.figure(figsize = (20, 8))
df.boxplot(rot = 45)
```

Out[16]: <Axes: >



### Check the total outliers of each columns

```
In [17]: total_outliers = []

for col in df.columns[1:]:
    percentile25 = df[col].quantile(0.25)
    percentile75 = df[col].quantile(0.75)
    iqr = percentile75 - percentile25
    upper_bound = percentile75 + 1.5 * iqr
    lower_bound = percentile25 - 1.5 * iqr
    total_outliers.append(sum((df[col] < lower_bound) | (df[col] > upper_bound)))
```

```
In [18]: total_outliers = pd.DataFrame(total_outliers)
```

```
In [19]: result_outlier = pd.concat([Col, total_outliers], axis = 1)
result_outlier.columns = ['Columns', 'Sum of Outliers']
```

In [20]:


result\_outlier

Out[20]:

	Columns	Sum of Outliers
0	CUST_ID	695.0
1	BALANCE	1493.0
2	BALANCE_FREQUENCY	808.0
3	PURCHASES	1013.0
4	ONEOFF_PURCHASES	867.0
5	INSTALLMENTS_PURCHASES	1030.0
6	CASH_ADVANCE	0.0
7	PURCHASES_FREQUENCY	782.0
8	ONEOFF_PURCHASES_FREQUENCY	0.0
9	PURCHASES_INSTALLMENTS_FREQUENCY	525.0
10	CASH_ADVANCE_FREQUENCY	804.0
11	CASH_ADVANCE_TRX	766.0
12	PURCHASES_TRX	248.0
13	CREDIT_LIMIT	808.0
14	PAYMENTS	774.0
15	MINIMUM_PAYMENTS	1474.0
16	PRC_FULL_PAYMENT	1366.0
17	TENURE	NaN

## Checking the Distribution of Data

```
In [ ]: for col in df.columns:
        sns.histplot(df[col], kde = True, bins = 20)
        plt.title('Distribution of {}'.format(col))
        plt.xlabel(col)
        plt.ylabel('FREQUENCY')
        plt.show()
```



## Statistical Analysis

```
In [21]: df.describe()
```

```
Out[21]:
```

	CUST_ID	BALANCE	BALANCE_FREQUENCY	PURCHASES	ONEOFF_PURCHASES	II
<b>count</b>	8950.000000	8950.000000	8950.000000	8950.000000	8950.000000	
<b>mean</b>	14600.040670	1564.474828	0.877271	1003.204834	592.437371	
<b>std</b>	2651.305875	2081.531879	0.236904	2136.634782	1659.887917	
<b>min</b>	10001.000000	0.000000	0.000000	0.000000	0.000000	
<b>25%</b>	12307.250000	128.281915	0.888889	39.635000	0.000000	
<b>50%</b>	14598.500000	873.385231	1.000000	361.280000	38.000000	
<b>75%</b>	16899.750000	2054.140036	1.000000	1110.130000	577.405000	
<b>max</b>	19190.000000	19043.138560	1.000000	49039.570000	40761.250000	

In [22]: *# Data Splitting*

```
X = df.iloc[:,:]
```

In [23]: *# Checking the number of Rows and Columns*  
X.shape

Out[23]: (8950, 18)

## Algorithm Selection



## K-Means with Elbow Method

```
In [24]: distances = []
from sklearn.cluster import KMeans
for i in range(1,11):
    kmeans = KMeans(n_clusters=i, init = 'k-means++', random_state = 42)
    kmeans.fit(X)
    distances.append(kmeans.inertia_)
```

/usr/local/lib/python3.10/dist-packages/sklearn/cluster/\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning


warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

warnings.warn(  
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/\_kmeans.py:870: FutureWarning: The default value of `n\_init` will change from 10 to 'auto' in 1.4. Set the value of `n\_init` explicitly to suppress the warning

**Plot distance with their respective number of clusters**

```
In [25]: plt.plot(range(1,11), distances)
plt.title('Elbow Method')
plt.xlabel('Clusters')
plt.ylabel('Distances')
plt.show()
```






```
In [34]: # Find suitable number of clusters with high silhouette score
distances
```

```
Out[34]: [0.42593449979764775,
0.3758402535764535,
0.25130717076663545,
0.2590742535484474,
0.2678799817209666,
0.2713220741254244,
0.2775015035460162,
0.2824372102444667,
0.28558068256430075]
```

### Plotting # of Clusters with their Silhouette Score

```
In [35]: plt.scatter(range(2,11), distances)
```

```
Out[35]: <matplotlib.collections.PathCollection at 0x7e4eccc6bf0>
```



### Model Implementation with KMeans Clustering

```
In [26]: k_cluster = KMeans(n_clusters=2, init = 'k-means++', random_state = 42)
elbow_pred = k_cluster.fit_predict(X)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_kmeans.py:870: FutureWarning: The default value of `n_init` will change from 10 to 'auto' in 1.4. Set the value of `n_init` explicitly to suppress the warning
warnings.warn(
```

## KMeans Clustering using Iterative Method

```
In [37]: # Converting DataFrame to Numpy array
```

```
X = np.array(X)
X.shape
```

```
Out[37]: (8950, 18)
```

### Find Closest Centroid

```
In [29]: def find_close_centroid(X, initial_centroids):
    idx = np.zeros(X.shape[0], 'float64')
    lst = []
    for i in range(X.shape[0]):
        lst = np.sum((X[i] - initial_centroids) ** 2, axis = 1)
        idx[i] = np.argmin(lst)

    return idx
```

### Compute Centroid to get Mean

```
In [30]: def compute_centroid(X, idx, K):
    m,n = X.shape
    centroids = np.zeros((K,n))
    for i in range(K):
        centroids[i] = np.mean(X[idx == i], axis = 0)

    return centroids
```

### Find the best Centroid using iteration

```
In [31]: def Run_KMean(X, initial_centroids,max_iteration = 10):
    idx = np.zeros(X.shape[0])
    centroids = initial_centroids.copy()
    K = initial_centroids.shape[0]
    for i in range(max_iteration):

        idx = find_close_centroid(X, initial_centroids)
        centroids = compute_centroid(X, idx, K)
    return idx, centroids
```

### Parameter Initialization

```
In [36]: k = 2
rand_data = np.random.permutation(X.shape[0])
initial_centroids = X[rand_data[:k]]
max_iter = 10
idx, centroids = Run_KMean(X, initial_centroids, max_iter)
idx
centroids
```


```
Out[36]: array([[1.46622152e+04, 6.95757523e+02, 8.46045865e-01, 6.60864763e+02,
3.50396867e+02, 3.10773477e+02, 3.20648284e+02, 4.94836389e-01,
1.75282869e-01, 3.69355034e-01, 8.23233760e-02, 1.60431428e+00,
1.14906314e+01, 3.21609479e+03, 9.46435686e+02, 4.22468591e+02,
1.80652394e-01, 1.14885845e+01],
[1.44481089e+04, 3.68730038e+03, 9.53572800e-01, 1.83975804e+03,
1.18389533e+03, 6.56149700e+02, 2.58732559e+03, 4.79388805e-01,
2.68862935e-01, 3.52420310e-01, 2.64218866e-01, 7.26741054e+00,
2.25763755e+01, 7.61827801e+03, 3.65556923e+03, 1.94365161e+03,
8.78887076e-02, 1.15875337e+01]])
```

## Hierarchical Clustering

## Plot Dendrogram

```
In [ ]: # Plotting Dendrogram for appropriate number of clusters

import scipy.cluster.hierarchy as sch
dendrogram = sch.dendrogram(sch.linkage(X, method = 'ward'))
plt.title('Dendrogram')
plt.xlabel('Customers')
plt.ylabel('Euclidean Distance')
plt.show()
```



## Model Implementation with Hierarchical Clustering

```
In [38]: from sklearn.cluster import AgglomerativeClustering
hc = AgglomerativeClustering(n_clusters = 2, affinity = 'euclidean', linkage = 'ward')
hc_pred = hc.fit_predict(X)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/cluster/_agglomerative.py:983:
FutureWarning: Attribute `affinity` was deprecated in version 1.2 and will be removed in 1.4. Use `metric` instead
  warnings.warn(
```

# Compute the Differences of each Model

```
In [39]: elbow_new = pd.DataFrame(elbow_pred)
         hc_new = pd.DataFrame(hc_pred)
         idx_new = pd.DataFrame(idx)
```

```
In [40]: table = pd.concat([elbow_new, idx_new, hc_new], axis = 1)
```

```
In [41]: table.columns = ['Elbow Cluster', 'IDX Cluster', 'Hierarchy Cluster',]
```

```
In [42]: table
```

Out[42]:

	Elbow Cluster	IDX Cluster	Hierarchy Cluster
0	0	0.0	1
1	1	1.0	0
2	1	0.0	0
3	1	0.0	0
4	0	0.0	1
...	...	...	...
8945	0	0.0	1
8946	0	0.0	1
8947	0	0.0	1
8948	0	0.0	1
8949	0	0.0	1

8950 rows × 3 columns

